

Non-Locality in NCA

Sreedev M

August 05, 2026

Contents

1. Motivation	3
1.1. Modelling Life as Interactions	3
1.2. Are neighbours enough?	3
1.3. Can NCA compute everything?	4
1.4. Why bother?	4
2. Experiments	5
2.1. Estimating Global Properties	5
2.1.1. Cell Differentiation Task	5
Bibliography	6

1. Motivation

1.1. Modelling Life as Interactions

How do biological systems, composed of numerous individual units, achieve many incredible feats such as coordination and stable complex pattern formation? In a 1952 paper titled *The Chemical Basis of Morphogenesis*, Alan Turing presents a mathematical model where the interaction between two diffusing substances giving rise to intricate patterns [1]. The diffusing substances, called “morphogens”, are modelled with the help linear differential equations. Differential equations are local, meaning the evolution of the system at a point is determined only by its neighbours, both in space and time. This choice is based on the assumption that the basic process in biological computation is a cell reacting to its local environment. Turing’s reaction diffusion model successfully modelled various complex patterns seen in nature, and it even got adopted to many other fields like physics, economics, sociology etc. The “local computations leading to global emergent complex behaviour” theory became an integral part of biological modelling henceforth.

Cellular Automata (CA) are models of computations introduced by John Von Neumann, where computation occurs via the interaction between “cells” in a grid. The cells can exist in a finite number of states. The next state of a cell is determined by its current state, and the state of its immediate neighbours. CAs are very powerful, they are even shown to be Turing Complete. G. Ermentrout and L. Edelstein-Keshet, employed CAs to model various developmental, populational biology patterns and a neuronal model. CA have been used in various biological settings, since its implicit local and simple update rules grants control to model things at a cellular level and let interactions among cells take care of the resultant complex emergent behaviour.

Neural Cellular Automata, is a modification on CA architecture. Instead of hand-coding the rules, a MLP is used which learns the local update rules of the CA, given a final state of the CA as target. NCA have been trained to model growth from a single cell, generate stable and self-stabilising patterns, and regenerate from a damaged state. Even more, they can self classify images, solve mazes and so on. One can think of NCA as machines that solve a task, with only means of localized computations. Efforts are underway now for establishing NCA as Universal Computing architectures.

1.2. Are neighbours enough?

Any information can spread within a grid, via cell to cell communication, **given enough time**. But there are many systems in nature, where we can find instances of non-local communication. Brain, nature’s most powerful and efficient computer is a good example. There are evidences of non-

local communications between different segments of the brain. In a much simpler system like the zebrafish skin, which has a characteristic gold and black pattern, there are long range connections called “airnemes” that help with the pattern stability. If only cell-to-cell communications are sufficient for information propagation, why did nature evolve such systems is a natural question to pose and the answer may also come as obvious. One is that diffusion is slow, and a system might have to function in conditions where information transfer via diffusion is too slow, which led to the formation of long-range jumps. Other one is that of efficiency, where the argument is that instead of wiring every cell with its neighbour, it might be more efficient to skip some of those wirings and add long range jumps. Another reason might be for the need to position a cell unit. In CAs, cells are symmetric; which implies they have no mechanism to gauge their position within the grid. But organisms have positional and body axis information, which might be an artefact emerging from long-range connections.

1.3. Can NCA compute everything?

Nature, having evolved non-local connections to facilitate such tasks implies that certain computations are either not feasible or not possible with NCA. While there are NCA with flexible topology called Graph-NCA, there have been little to no studies comparing vanilla NCA to GNCA. Especially, there have been very no studies gauging what the limitations of NCA are and what additional capabilities non-local connections impart to them. To gauge the computational bounds or limitations of NCA, one can cook up toy problems of scenarios where nature has evolved non-local connections and assess the performance of NCA on them. The intuition is that vanilla NCA will perform substantially poorly than a NCA with long range connections.

1.4. Why bother?

One straightforward reason is to assess what NCA can compute and cannot/struggle to compute. This can also serve as a comparison of vanilla NCA to GNCA approaches, and a study of the capabilities non-locality imparts to the architecture exclusively. Another more ambitious aspiration, is that if the experiments turn out to be fruitful, a non-local NCA can serve as a tool to examine the reasons for emergence of non-local systems in nature and life. Testing new hypotheses that try to explain why non-local structures emerged, either functionally or evolutionarily can be tested with a toy problem here. We can vary the topology and connection architecture to see if feasible alternative pathways are allowed in nature.

- meaning it might take a huge amount of time for the NCA to compute
- or GNCA for short
- with Moore neighbours; the immediate square neighbourhood of 8
- The nature and structure of those connections can be taken up depending upon the nature of the problem.

2. Experiments

2.1. Estimating Global Properties

Usually, NCA are trained on tasks such that each cell has a predefined target to reach. What the network learns is to how to guide each cell, only using it's current value the value of its current neighbors, to that target state. NCA shows remarkable performance in such tasks.

- something like generating an image, where each cell has to reach it's target pixel

The goal of the experiments in this section is not to have a specified target for each cell. The loss is predominantly applicable for the whole grid, and the values taken by each individual cell is of less importance than the global properties of the grid that emerge from those values. The intuition here is to test how NCA fares in situations where each cell has to perform by taking into account how the global property of the grid might change due to it's actions.

- this is not exactly true, but I'll clarify this later.

2.1.1. Cell Differentiation Task

The task here is described as given below:

1. Each grid starts out

Bibliography

- [1] A. M. Turing, “The chemical basis of morphogenesis,” *Philosophical Transactions of the Royal Society of London. B, Biological Sciences*, vol. 237, no. 641, pp. 37–72, 1952, doi: 10.1098/rstb.1952.0012[◦].
- [2] G. Ermentrout and L. Edelstein-Keshet, “Cellular Automata Approaches to Biological Modeling,” *Journal of Theoretical Biology*, vol. 160, no. 1, pp. 97–133, 1993, doi: <https://doi.org/10.1006/jtbi.1993.1007>[◦].